

Midterm Project

Ekta Chaudhary

26/03/2020

```
library(tidyverse)
library(readxl)
library(caret)
library(ModelMetrics)
library(glmnet)
library(gam)
library(mgcv)
library(splines)
library(pdp)
library(earth)
library(dplyr)
library(naniar)
library(bnstruct)
library(corrplot)
library(logisticPCA)
library(MASS)
library(e1071)
library(mlbench)
library(pROC)
library(AppliedPredictiveModeling)
```

Reading the Datasets. Adding the column names to the Breast_Cancer dataset.

```
Diagnosis =
read_csv('./data/Diagnosis.csv')
Prognosis =
  read_csv('./data/Prognosis.csv')
Breast_Cancer =
  read_csv(file = './data/Breast_Cancer.csv' ,col_names = c('id_number','Clump_thickness','Uniformity_o
```

Replacing the missing observations that are denoted by a ? with na and then using KNN imputation to impute the missing values.

```
Breast_Cancer = Breast_Cancer %>% replace_with_na_all(condition = ~.x == '?')
```

```
Breast_Cancer <- knn.impute(as.matrix(Breast_Cancer), k = 10, cat.var = 2:ncol(Breast_Cancer) - 2,
  to.impute = 1:nrow(Breast_Cancer), using = 1:nrow(Breast_Cancer))
```

Creating a data frame called Breast_Cancer.


```
Breast_Cancer <- data.frame(Breast_Cancer)
```

Using ifelse on the column Class_cancer -> If Class_cancer = 4 then it's malignant, else its benign. Converted the Class_cancer into factors.

```
Breast_Cancer$Class_cancer = as.factor(ifelse(Breast_Cancer$Class_cancer == 4, 'mal', 'benign'))
Breast_Cancer = Breast_Cancer[,2:11]
```

Creating a correlation plot to check the correlation between the variables.

```
x = model.matrix(Class_cancer~., Breast_Cancer) [, -1]
corrplot(cor(x))
```



Here we are diving the data into training and test data.

```
set.seed(1)
rowTrain <- createDataPartition(y = Breast_Cancer$Class_cancer,
                                p = 2/3,
                                list = FALSE)
```

Logistic Regression:

```
glm.fit <- glm(Class_cancer~.,
               data = Breast_Cancer,
               subset = rowTrain,
               family = binomial)

contrasts(Breast_Cancer$Class_cancer)
```


```
##          mal
```

```
## benign    0
## mal       1
```

We first consider the Bayes classifier (cutoff 0.5) and evaluate its performance on the test data and then plot the test ROC curve.

```
test.pred.prob <- predict(glm.fit, newdata = Breast_Cancer[-rowTrain,1:9],
                          type = "response")
test.pred <- rep("benign", length(test.pred.prob))
test.pred[test.pred.prob > 0.5] <- "mal"

roc.glm <- roc(Breast_Cancer$Class_cancer[-rowTrain], test.pred.prob)
plot(roc.glm, legacy.axes = TRUE, print.auc = TRUE)
plot(smooth(roc.glm), col = 4, add = TRUE)
```



We can also fit a logistic regression using caret. This is to compare the cross-validation performance with other models, rather than tuning the model.

```
ctrl <- trainControl(method = "repeatedcv",
                    repeats = 5,
                    summaryFunction = twoClassSummary,
                    classProbs = TRUE)
```

```
set.seed(1)
model.glm <- train(x = Breast_Cancer[rowTrain,1:9],
                  y = Breast_Cancer$Class_cancer[rowTrain],
                  method = "glm",
                  metric = "ROC",
                  trControl = ctrl)
```

```
coef(glm.fit)
```

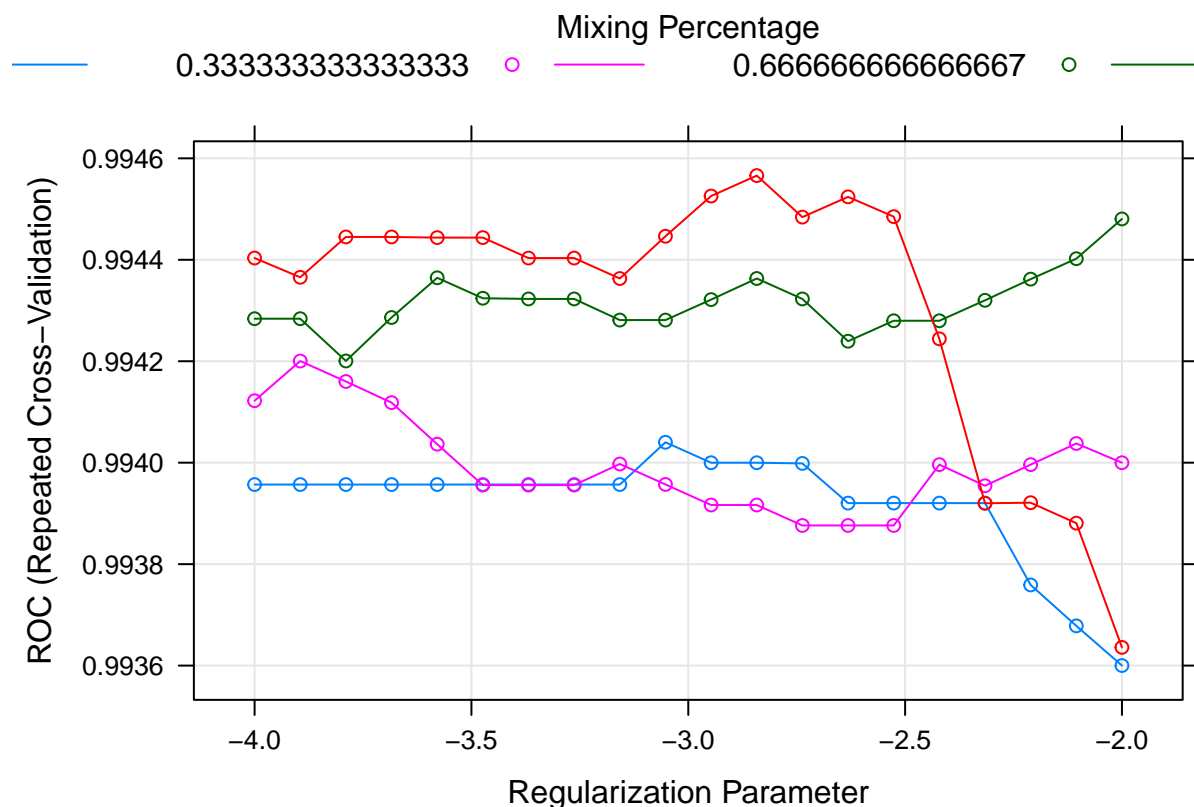
```
##           (Intercept)           Clump_thickness  Uniformity_of_Cell_Size
##           -9.0107637           0.4221055           0.1457212
## Uniformity_of_Cell_Shape  Marginal_Adhesion           Epithelial_Cell_Size
##           0.6370116           0.1479078           -0.2071035
##           Bare_Nuclei           Bland_Chromatin           Normal_Nucleoli
##           0.5172205           0.2023960           0.1863717
##           Mitoses
##           0.6735681
```

Regularized logistic regression can be fitted using `glmnet`. We use `thetrain` function to select the optimal tuning parameters.

```
glmGrid <- expand.grid(.alpha = seq(0, 1, length = 4),
                      .lambda = exp(seq(-4, -2, length = 20)))
```

```
set.seed(1)
model.glmn <- train(x = Breast_Cancer[rowTrain,1:9],
                    y = Breast_Cancer$Class_cancer[rowTrain],
                    method = "glmnet",
                    tuneGrid = glmGrid,
                    metric = "ROC",
                    trControl = ctrl)
```

```
plot(model.glmn, xTrans = function(x) log(x))
```





```
model.glmn$bestTune
```

```
##      alpha      lambda  
## 72      1 0.05830279
```

```
test.pred.prob <- predict(model.glmn, x = Breast_Cancer[-rowTrain,1:9],  
                           method = "glmnet")
```

Here we use the function `lda` in library `MASS` to conduct LDA.

```
library(MASS)  
  
lda.fit <- lda(Class_cancer~., data = Breast_Cancer,  
              subset = rowTrain)  
plot(lda.fit)
```




Here we are evaluating the test set performance using ROC.

```
lda.pred <- predict(lda.fit, newdata = Breast_Cancer[-rowTrain,])  
head(lda.pred$posterior)
```

```
##      benign      mal  
## 3 0.99998228 1.772454e-05  
## 4 0.00875467 9.912453e-01  
## 5 0.99998862 1.137890e-05  
## 7 0.81869551 1.813045e-01  
## 8 0.99999630 3.695724e-06  
## 9 0.99999902 9.816974e-07
```

```
roc.lda <- roc(Breast_Cancer$Class_cancer[-rowTrain], lda.pred$posterior[,2],
              levels = c("benign", "mal"))

plot(roc.lda, legacy.axes = TRUE, print.auc = TRUE)
```




Using caret:

```
set.seed(1)
model.lda <- train(x = Breast_Cancer[rowTrain,1:9],
                  y = Breast_Cancer$Class_cancer[rowTrain],
                  method = "lda",
                  metric = "ROC",
                  trControl = ctrl)
```

Using KNN:

```
set.seed(1)
model.knn <- train(x = Breast_Cancer[rowTrain,1:9],
                  y = Breast_Cancer$Class_cancer[rowTrain],
                  method = "knn",
                  preProcess = c("center", "scale"),
                  tuneGrid = data.frame(k = seq(1,50,by = 5)),
                  trControl = ctrl)

ggplot(model.knn)
```



GLM, Regularized GLM and LDA have relatively good performance.

```
res <- resamples(list(GLM = model.glm, GLMNET = model.glmn,
                      LDA = model.lda, KNN = model.knn))
summary(res)
```

```
##
## Call:
## summary.resamples(object = res)
##
## Models: GLM, GLMNET, LDA, KNN
## Number of resamples: 50
##
## ROC
##           Min.   1st Qu.   Median     Mean 3rd Qu.  Max. NA's
## GLM      0.9556452 0.9883233 0.9979503 0.9935506      1      1      0
## GLMNET   0.9576613 0.9917339 0.9979839 0.9945660      1      1      0
## LDA      0.9556452 0.9905637 0.9979839 0.9941979      1      1      0
## KNN      0.9506048 0.9881048 0.9964767 0.9910701      1      1      0
##
## Sens
##           Min.   1st Qu.   Median     Mean 3rd Qu.  Max. NA's
## GLM      0.9333333 0.9666667 0.983871 0.9777419      1      1      0
## GLMNET   0.9354839 0.9677419 1.000000 0.9875699      1      1      0
## LDA      0.9333333 0.9677419 1.000000 0.9855914      1      1      0
## KNN      0.9333333 0.9666667 0.983871 0.9803656      1      1      0
##
```

```
## Spec
##           Min.      1st Qu.      Median      Mean      3rd Qu.      Max.      NA's
## GLM      0.8125 0.8750000 0.9375000 0.9363235 1.0000000      1      0
## GLMNET   0.6875 0.8235294 0.8750000 0.8942647 0.9375000      1      0
## LDA      0.6875 0.8750000 0.8786765 0.9003676 0.9375000      1      0
## KNN      0.7500 0.8750000 0.9375000 0.9188971 0.9852941      1      0
```

Now looking at the test set performance.

```
lda.pred <- predict(model.lda, newdata = Breast_Cancer[-rowTrain,], type = "prob")[,2]
glm.pred <- predict(model.glm, newdata = Breast_Cancer[-rowTrain,], type = "prob")[,2]
glmnpred <- predict(model.glmn, newdata = Breast_Cancer[-rowTrain,], type = "prob")[,2]
knn.pred <- predict(model.knn, newdata = Breast_Cancer[-rowTrain,], type = "prob")[,2]

roc.lda <- roc(Breast_Cancer$Class_cancer[-rowTrain], lda.pred)
roc.glm <- roc(Breast_Cancer$Class_cancer[-rowTrain], glm.pred)
roc.glmn <- roc(Breast_Cancer$Class_cancer[-rowTrain], glmnpred)
roc.knn <- roc(Breast_Cancer$Class_cancer[-rowTrain], knn.pred)

auc <- c(roc.glm$auc[1], roc.glmn$auc[1], roc.lda$auc[1], roc.knn$auc[1])

plot(roc.glm, legacy.axes = TRUE)
plot(roc.glmn, col = 2, add = TRUE)
plot(roc.lda, col = 3, add = TRUE)
plot(roc.knn, col = 6, add = TRUE)
modelNames <- c("glm", "glmn", "lda", "knn")
legend("bottomright", legend = paste0(modelNames, ": ", round(auc,3)),
      col = 1:6, lwd = 2)
```

